

计算机系统导论期中考试 2012 年 11 月 7 日

姓名:

学号:

成绩:

Problem A -- Multiple Choice (18 pts, 1pt or 2pts each)

1. Let `int x = -31/8` and `int y = -31 >> 3`. What are the values of `x` and `y`? (m-f10-a 1-5) C
 - a) `x = -3, y = -3`
 - b) `x = -4, y = -4`
 - c) `x = -3, y = -4`
 - d) `x = -4, y = -3`
2. For which values can `X` not be equal to `Z` in the code below (circle all that apply) (f-f09 1-B) A&D

```
int X = CONSTANT;
float Y = X;
int Z = Y;
```

 - a) For large positive values of `CONSTANT` (e.g., >1,000,000,000)
 - b) For large negative values of `CONSTANT` (e.g., >-100)
 - c) For small positive values of `CONSTANT` (e.g., <100)
 - d) For small negative values of `CONSTANT` (e.g., <-1,000,000,000)
 - e) None of the above (i.e., `X==Z` in all of these cases)
3. Consider the following code, what is the output of the `printf`? (m-f10-a 1-1) B

```
int x = 0x15213F10 >> 4;
char y = (char) x;
unsigned char z = (unsigned char) x;
printf("%d, %u", y, z);
```

 - a) -241, 15
 - b) -15, 241
 - c) -241, 241
 - d) -15, 15
4. After executing the following code, which of the variables are equal to 0? (m-f11 1-3) CD or CDE
 - a) `unsigned int a = 0xffffffff;`
 - b) `unsigned int b = 1;`
 - c) `unsigned int c = a + b;`
 - d) `unsigned long d = a + b;`
 - e) `unsigned long e = (unsigned long)a + b`
5. How does x86 assembly store the return value when a function is finished? (f-f10 1-4) B
 - a) The `ret` instruction stores it in a special `retval` register.
 - b) By convention, it is always in `%eax`.
 - c) It is stored on the stack just above the `(%ebp)` of the callee.
 - d) It is stored on the stack just above all the arguments to the function
6. Which of the following is FALSE concerning x86-64 architecture? (f-f10 1-9) D
 - a) A double is 64 bits long.
 - b) Registers are 64 bits long.

- c) Pointers are 64 bits long.
 - d) Pointers point to locations in memory that are multiples of 64 bits apart.
7. Denormalized floating point numbers are **(m-f10-a 1-18) A**
- a) Very close to zero (small magnitude)
 - b) Very far from zero (large magnitude)
 - c) Un-representable on a number line
 - d) Zero.
8. What is the minimum (most negative) value of a 32-bit two's complement integer? **(m-f11 1-1) C**
- a) -2^{32}
 - b) $-2^{32} + 1$
 - c) -2^{31}
 - d) $-2^{31} + 1$
9. What is the difference between the mov and lea instructions? **(m-f11 1-2) B**
- a) lea dereferences an address, while mov doesn't.
 - b) mov dereferences an address, while lea doesn't.
 - c) lea can be used to copy a register into another register, while mov cannot.
 - d) mov can be used to copy a register into another register, while lea cannot.
10. Which of the following 8 bit floating point numbers (1 sign, 3 exponent, 4 fraction) represent NaN? **(m-s11 1-2) B**
- a) 1 000 1111
 - b) 0 111 1111
 - c) 0 100 0000
 - d) 1 111 0000
11. `%rsp` is 0xdeadbeefdeadd0d0. What is the value in `%rsp` after the following instruction executes?
`pushq %rbx` **(m-s11 1-3) D**
- a) 0xdeadbeefdeadd0d4
 - b) 0xdeadbeefdeadd0d8
 - c) 0xdeadbeefdeadd0cc
 - d) 0xdeadbeefdeadd0c8
12. Consider an `int *a` and an `int n`. If the value of `%ecx` is `a` and the value of `%edx` is `n`, which of the following assembly snippets best corresponds to the C statement `return a[n]`? **(m-s11 1-1) C**
- a) `ret (%ecx,%edx,4)`
 - b) `leal (%ecx,%edx,4),%eax`
`ret`
 - c) `mov (%ecx,%edx,4),%eax`
`ret`
 - d) `mov (%ecx,%edx,1),%eax`
`ret`
13. What is the C equivalent of `mov 0x10 (%rax,%rcx,4), %rdx` **(2pts) (m-s11 1-8) C**
- a) `rdx = rax + rcx + 4 + 10`
 - b) `*(rax + rcx + 4 + 10) = rdx`
 - c) `rdx = *(rax + rcx*4 + 0x10)`
 - d) `rdx = *(rax + rcx + 4 + 0x10)`

14. In what section of an ELF binary are initialized variables located? **(2pts) (e2b-s11) B**
- a) .symtab
 - b) .data
 - c) .bss
 - d) .text
15. A system uses a two-way set-associative cache with 16 sets and 64-byte blocks. Which set does the byte with the address 0xdeadbeef map to? **(2pts) (e2-f10) B**
- a) Set 7
 - b) Set 11
 - c) Set 13
 - d) Set 14

Problem B -- Bits & Bytes, Int, and FP (20 pts)

1. For each of the following propositions, write in all comparisons that make it always true among the four possibilities:

`< > == !=`

If none are guaranteed to hold, please indicate that explicitly by marking it with an X. We have filled in the first two for you as examples. Assume the variables are declared with

`int x, y;`

and initialized to some unknown values. You may assume that int's are 32 bits wide, char's are 8 bits wide and that right shift is arithmetical on signed numbers and logical on unsigned numbers.

(m-s07 1) (4pts)

(1) !=

(2) <

!=

(3) >

!=

(4) !=

If $x > y$ then y	$\left\{ \begin{array}{c} < \\ != \end{array} \right\}$	x
If $x < 0$ then $x+1$	$\left\{ \begin{array}{c} \mathbf{X} \end{array} \right\}$	0
If $x > 0$ then $x+1$	$\left\{ \begin{array}{c} \end{array} \right\}$	0
If $(x \gg 31) < 0$ then x	$\left\{ \begin{array}{c} \end{array} \right\}$	0
If $((x \ll 31) \gg 31) < 0$ then $x \& 1$	$\left\{ \begin{array}{c} \end{array} \right\}$	0
If $x > y$ then $(\text{unsigned})x$	$\left\{ \begin{array}{c} \end{array} \right\}$	y

2. Floating point encoding. Consider the following 5-bit floating point representation based on the IEEE floating point format. This format does not have a sign bit – it can only represent nonnegative numbers.

- There are $k = 3$ exponent bits. The exponent bias is 3.
- There are $n = 2$ fraction bits.

Recall that numeric values are encoded as a value of the form $V = M \times 2^E$, where E is the exponent after biasing, and M is the significant value. The fraction bits encode the significant value M using either a denormalized (exponent field 0) or a normalized representation (exponent field nonzero). The exponent E is given by $E = 1 - \text{Bias}$ for denormalized values and $E = e - \text{Bias}$ for normalized values, where e is the value of the exponent field exp interpreted as an unsigned number.

Below, you are given some decimal values, and your task is to encode them in floating point format. In addition, you should give the rounded value of the encoded floating point number. To get credit, you must give these as whole numbers (e.g., 17) or as fractions in reduced form (e.g., $3/4$). Any rounding of the significand is based on round-to-even, which rounds an unrepresentable value that lies halfway between two representable values to the nearest even representable value.

(8pts)

Value	FP bits	Rounded value
3	100 10	3 normalized, exact
9	110 00	8 normalized, round down, because 10 is odd
3/16	000 11	3/16 exact, denorm
15/2	110 00	8 normalized, round up, change exponent

Value	Floating Point Bits	Rounded value
9/32	001 00	1/4
3		
9		
3/16		
15/2		

3. Conversion from float to int is a notoriously expensive operation. Not all processors do this in hardware. A clever technique uses the normalization step in double-precision floats. Imagine if you could normalize a (non-negative) floating point number so that the low-order significant bit represents 2^0 . Then the significant would be an integer. **(8 pts) (f-f09 7)**

Assume a double where, after taking the exponent into account, the low-order bit represents 1 (2^0). What does the “implied leading 1” of the significant represent?

2^{52}

Now, assume a double x is known to be in the range $0 \leq x < 232$. To force the low-order bit of x to represent 1 (2^0), you should add by **2^{52}** ?

After this operation, you can store the double to memory and read the (circle one:)

- Upper 32 bits (the end with the sign bit == 0), or
- **Lower 32 bits (the other end)**

as an unsigned integer to get the integer value of the original double. This algorithm will round (circle one)

- toward zero,
- up,
- down,
- **to the nearest even.**

Hints:

double-precision significant has 52 bits

double-precision exponent has 11 bits

double-precision sign-bit is (of course) 1 bit

Problem C Machine Programming (25 pts)

1. The function below is hand-written assembly code for a sorting algorithm. Fill in the blanks on the next page by converting this assembly to C code **(m-s09-4) (13pts)**

```
.globl mystery_sort      # exports the symbol so other .c files
                          # can call the function

mystery_sort:
    jmp     loop1_check

loop1:
    xor     %rdx, %rdx
    mov     %rsi, %rcx
    jmp     loop2_check

loop2:
    mov     (%rdi, %rcx, 8), %rax
    cmp     %rax, (%rdi, %rdx, 8)
    jg      loop2_check
    mov     %rcx, %rdx

loop2_check:
    dec     %rcx
    test    %rcx, %rcx
    jnz     loop2

    dec     %rsi
    mov     (%rdi, %rsi, 8), %rax
    mov     (%rdi, %rdx, 8), %rcx
    mov     %rcx, (%rdi, %rsi, 8)
    mov     %rax, (%rdi, %rdx, 8)

loop1_check:
    test    %rsi, %rsi
    jnz     loop1

    ret
```

```
void mystery_sort (long* array, long len)
{
    long a, b, tmp;
    while(len>0)
    {
        a = 0;
        for(b=len-1; b>0; b--)
        {
            if(array[b]>array[a])
            {
                a=b;
            }
        }
        len--;
        tmp = array[len];
        array[len] = array[a];
        array[a] = tmp;
    }
}
```

```

void mystery_sort (long* array, long len)
{
    long a, b, tmp;

    while (____ > ____ )
    {
        a = ____;

        for (b = ____; b > ____; b--)
        {
            if (array[____] > array{____])
            {
                ____ = ____;
            }
        }

        len--;

        tmp = array[____];
        array[____] = array[____];
        array[____] = tmp;
    }
}

```

2. Below is some assembly code to a famous algorithm. Please briefly read the code then answer the questions on the following page. **(m-s09-6) 12pts**

```

0000000000400498 <mystery>:
400498: 41 b8 00 00 00 00  mov    $0x0,%r8d
40049e: eb 22              jmp    4004c2 <mystery+0x2a>
4004a0: 89 c8              mov    %ecx,%eax
4004a2: c1 e8 1f          shr    $0x1f,%eax
4004a5: 01 c8              add    %ecx,%eax
4004a7: d1 f8              sar    %eax          ; arith shift right 1 bit
4004a9: 42 8d 0c 00       lea    (%rax,%r8,1),%ecx
4004ad: 48 63 c1          movslq %ecx,%rax
4004b0: 8b 04 87          mov    (%rdi,%rax,4),%eax
4004b3: 39 d0              cmp    %edx,%eax
4004b5: 7d 05              jge    4004bc <mystery+0x24>
4004b7: 41 89 c8          mov    %ecx,%r8d
4004ba: eb 06              jmp    4004c2 <mystery+0x2a>
4004bc: 39 d0              cmp    %edx,%eax
4004be: 7e 10              jle    4004d0 <mystery+0x38>
4004c0: 89 ce              mov    %ecx,%esi
4004c2: 89 f1              mov    %esi,%ecx
4004c4: 44 29 c1          sub    %r8d,%ecx
4004c7: 85 c9              test   %ecx,%ecx
4004c9: 7f d5              jg     4004a0 <mystery+0x8>
4004cb: b9 ff ff ff ff    mov    $0xffffffff,%ecx
4004d0: 89 c8              mov    %ecx,%eax
4004d2: c3                retq

```

- a) Please write a single line of C code to represent the instruction `lea (%rax,%r8,1),%ecx` (Use C variables named `rax`, `r8`, and `ecx`, you can ignore types).
`ecx = rax + r8`
- b) Please write a single line of C code to represent the instruction `mov (%rdi,%rax,4),%eax` (Use C variables named `rdi` and `rax`, you can ignore types).
`eax = *(rdi + rax) // rdi is int *`
`eax = *(rdi + 4 * rax) // rdi is char *`
- c) Commonly found in assembly is the leave instruction; why is that instruction not in this code?
`no stack to return from, no push %ebp`
- d) You learned about two different architectures in class, IA32 and x86 64. What architecture is this code written for and what major downside would occur from using the other architecture
`x86_64`
`32-bit has less registers than 64-bit mode, so local variables would need to be stored on the stack. Stack accessing takes extra instructions and extra time.`

Problem D Stack (10 pts)

Stack discipline. Consider the following C code and assembly code for a recursive function . **(m-f11-5) 10pts**


```

int gcd(int a, int b)      0x08048394 <+0>:    push    %ebp
{                          0x08048395 <+1>:    mov     %esp, %ebp
    if(!b)                0x08048397 <+3>:    sub     $0x10, %esp
    {                      0x0804839a <+6>:    mov     0x8(%ebp), %eax
        return a;          0x0804839d <+9>:    mov     0xc(%ebp), %ecx
    }                     0x080483a0 <+12>:   test    %ecx, %ecx
                          0x080483a2 <+14>:   je      0x80483b7 <gcd+35>
    return gcd(b, a % b);  0x080483a4 <+16>:   mov     %eax, %edx
                          0x080483a6 <+18>:   sar     $0x1f, %edx
                          0x080483a9 <+21>:   idiv    %ecx
                          0x080483ab <+23>:   mov     %edx, 0x4(%esp)
                          0x080483af <+27>:   mov     %ecx, (%esp)
                          0x080483b2 <+30>:   call    0x8048394 <gcd>
                          0x080483b7 <+35>:   leave
                          0x080483b8 <+36>:   ret
    
```

Imagine that a program makes the procedure call `gcd(213, 18)`. Also imagine that prior to the invocation, the value of `%esp` is `0xffff1000`—that is, `0xffff1000` is the value of `%esp` immediately before the execution of the call instruction.

A. Note that the call `gcd(213, 18)` will result in the following function invocations: `gcd(213, 18)`, `gcd(18, 15)`, `gcd(15, 3)`, and `gcd(3, 0)`. Using the provided code and your knowledge of IA32 stack discipline, fill in the stack diagram with the values that would be present immediately before the execution of the `leave` instruction for `gcd(15, 3)`. Supply numerical values wherever possible, and cross out each blank for which there is insufficient information to complete with a numerical value.

Hints: The following set of C style statements describes an approximation of the operation of the instruction `idiv %ecx`, where `/` is the division operator and `%` is the modulo operator:

`%eax = %eax / %ecx`

`%edx = %eax % %ecx`

Also, recall that `leave` is equivalent to `movl %ebp, %esp; popl %ebp`

Part A:

+	-----+	
	unknown	0xffff1008
+	-----+	
	18	0xffff1004
+	-----+	
	213	0xffff1000
+	-----+	
	unknown	0xffff0ffc
+	-----+	
	unknown	0xffff0ff8
+	-----+	
	unknown	0xffff0ff4

unknown	0xffff0ff0
15	0xffff0fec
18	0xffff0fe8
0x080483b7	0xffff0fe4
0xffff0ff8	0xffff0fe0
unknown	0xffff0fdc
unknown	0xffff0fd8
3	0xffff0fd4
15	0xffff0fd0
0x080483b7	0xffff0fcc
0xffff0fe0	0xffff0fc8
unknown	0xffff0fc4
unknown	0xffff0fc0
0	0xffff0fbc
3	0xffff0fb8
unknown	0xffff0fb4
unknown	0xffff0fb0

B. What are the values of %esp and %ebp immediately before the execution of the ret instruction for gcd(15, 3)?

esp: 0xffff0fcc
ebp: 0xffff0fe0

Problem E -- The Hit or Miss Question (15 pts)

Given a 32-bit Linux system that has a 2-way associative cache of size 128 bytes with 32 bytes per block. Long longs are 8 bytes. For all parts, assume that table starts at address 0x0.

```
int i;
```

```
int j;
```

```
long long table[4][8];
```

```
for (j = 0; j < 8; j++) {
```

```
    for (i = 0; i < 4; i++) {
```

```
        table[i][j] = i + j;
```

```
    }
```

```
}
```

- A. This problem refers to code sample 1. In the table below write down in each space whether that element's access will be a hit or a miss. Indicate hits with a 'H' and misses with a 'M'.

	0	1	2	3	4	5	6	7
0	M	M	M	M	M	M	M	M
1	M	M	M	M	M	M	M	M
2	M	M	M	M	M	M	M	M
3	M	M	M	M	M	M	M	M

What is the miss rate of this code sample?

1

```
int i;

int j;

int table[4][8];

for (j = 0; j < 8; j++) {
    for (i = 0; i < 4; i++) {
        table[i][j] = i + j;
        table[j][i]?
    }
}
```

- B. This problem refers to code sample above. In the table below write down in each space whether that element's access will be a hit or a miss. Indicate hits with a 'H' and misses with a 'M'.

	0	1	2	3	4	5	6	7
0	M	H	H	H	H	H	H	H
1	M	H	H	H	H	H	H	H
2	M	H	H	H	H	H	H	H
3	M	H	H	H	H	H	H	H

What is the miss rate of this code sample?

1/8

- C. One code sample performs better than the other. Why is this?

The second sample performs better than the first because it makes better use of the cache. In the first, heavy conflict results in a large number of evictions. In the second, the entire array fits in the cache, and the only misses are cold misses.

Problem F – Linking (12 pts)

For each of the following code snippets, write down all symbols in the resulting object files from compilation. Write whether it is a weak global, strong global, or local variable, and what section of the final compiled ELF binary the variable will go into. Fill in the value if you have enough information to determine the value.

A.

main.c	foo.c
<pre>#include <stdio.h> int x; int y; int z = 0; int main() { printf("%x\n", x); printf("%x\n", y); x = 0xdeadbeef; printf("%x\n", x); printf("%x\n", y); return 0; }</pre>	<pre>short x = 5; short y = 2;</pre>

File	Symbol	Strength / scope	Value	ELF Section
Main.o	X	Weak global	-	.data
	Y	Weak global	-	.data
	Z	Strong global	0	.data
	Main	Strong global	-	.text
foo.o	X	Strong global	5	.data
	Y	Strong global	2	.data